

Hands-on Lab: Improving Performance of Slow Queries in MySQL

Estimated time needed: 45 minutes

In this lab, you will learn how to improve the performance of your slow queries in MySQL, which can be particularly helpful with large databases.

Objectives

After completing this lab, you will be able to:

1. Use the `EXPLAIN` statement to check the performance of your query
2. Add indexes to improve the performance of your query
3. Apply other best practices such as using the `USE INDEX` clause to improve query performance

Software Used in this Lab

In this lab, you will use [MySQL](#). MySQL is a Relational Database Management System (RDBMS) designed to efficiently store, manipulate, and retrieve data.




To complete this lab, you will utilize the MySQL relational database service available as part of the IBM Skills Network Labs (SN Labs) Cloud IDE. SN Labs is a virtual lab environment used in this course.

Database Used in this Lab

The Employees database used in this lab comes from the following source: <https://dev.mysql.com/doc/employees/en/> under the [CC BY-SA 3.0 license](#).

The following entity relationship diagram (ERD) shows the schema of the Employees database:



The first row of each table is the table name, the rows with keys next to them indicate the primary keys, and the remaining rows are additional attributes.

Exercise 1: Load the Database

Let's begin by retrieving the database and loading it so that it can be used.

1. In the menu bar, select `Terminal` > `New Terminal`. This will open the Terminal.

To download the zip file containing the database, copy and paste the following into the Terminal:

```
1 1
2 wget https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/IBM-DB0231EN-SkillsNetwork/datasets/employeesdb.zip
3
theia@theiadocker-: /home/project$ wget https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/IBM-DB0231EN-SkillsNetwork/datasets/employeesdb.zip
--2021-10-12 20:08:23-- https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/IBM-DB0231EN-SkillsNetwork/datasets/employeesdb.zip
Resolving cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud (cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud)... 198.23.119.245
Connecting to cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud (cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud)|198.23.119.245|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 36689578 (35M) [application/zip]
Saving to: 'employeesdb.zip'

employeesdb.zip      100% [=====] 34.99M  30.3MB/s   in 1.2s

2021-10-12 20:08:25 (30.3 MB/s) - 'employeesdb.zip' saved [36689578/36689578]

theia@theiadocker-: /home/project$
```

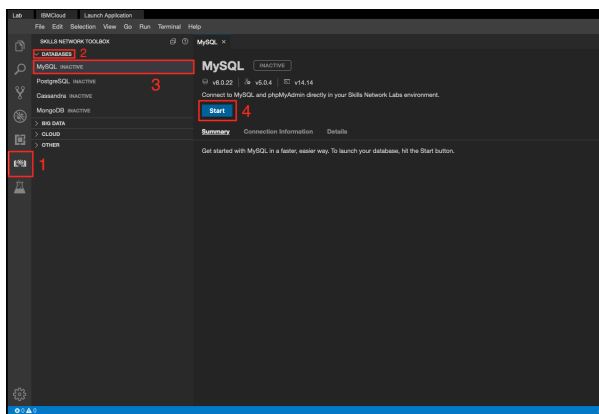
2. Next, we'll need to unzip its contents. We can do that with the following command:

```
1 1
2 unzip employeesdb.zip
3
theia@theiadocker-: /home/project$ unzip employeesdb.zip
Archive: employeesdb.zip
  creating: employeesdb/
  creating: employeesdb/sakila/
  inflating: employeesdb/load_salaries2.dump
  inflating: employeesdb/test_versions.sh
  inflating: employeesdb/objects.sql
  inflating: employeesdb/load_salaries3.dump
  inflating: employeesdb/load_dept_emp.dump
  inflating: employeesdb/test_employees_sha.sql
  inflating: employeesdb/ChangeLog
  creating: employeesdb/images/
  inflating: employeesdb/employees_partitioned_5.1.sql
  inflating: employeesdb/test_employees_md5.sql
  inflating: employeesdb/README.md
  inflating: employeesdb/employees.sql
  inflating: employeesdb/load_titles.dump
  inflating: employeesdb/employees_partitioned.sql
  inflating: employeesdb/load_dept_manager.dump
  inflating: employeesdb/sql_test.sh
  inflating: employeesdb/load_departments.dump
  inflating: employeesdb/load_salaries1.dump
  inflating: employeesdb/show_elapsed.sql
  inflating: employeesdb/load_employees.dump
  inflating: employeesdb/sakila/README.md
  inflating: employeesdb/sakila/sakila-mv-data.sql
  inflating: employeesdb/sakila/sakila-mv-schema.sql
  inflating: employeesdb/images/employees.jpg
  inflating: employeesdb/images/employees.png
  inflating: employeesdb/images/employees.gif
theia@theiadocker-: /home/project$
```

3. Now, let's change directories so that we're able to access the files in the newly created **employeesdb** folder.

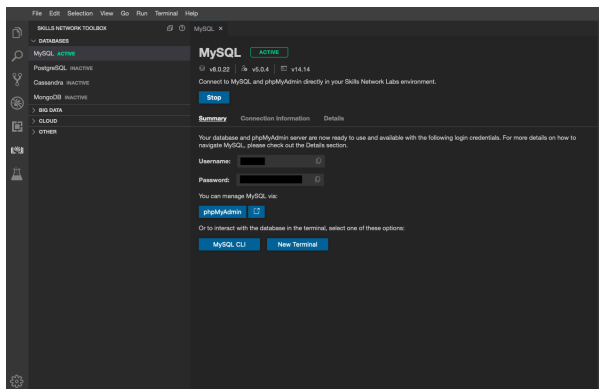
```
1 1
2 cd employeesdb
3
theia@theiadocker-: /home/project$ cd employeesdb
Archive: employeesdb.zip
  creating: employeesdb/
  creating: employeesdb/sakila/
  inflating: employeesdb/load_salaries2.dump
  inflating: employeesdb/test_versions.sh
  inflating: employeesdb/objects.sql
  inflating: employeesdb/load_salaries3.dump
  inflating: employeesdb/load_dept_emp.dump
  inflating: employeesdb/test_employees_sha.sql
  inflating: employeesdb/ChangeLog
  creating: employeesdb/images/
  inflating: employeesdb/employees_partitioned_5.1.sql
  inflating: employeesdb/test_employees_md5.sql
  inflating: employeesdb/README.md
  inflating: employeesdb/employees.sql
  inflating: employeesdb/load_titles.dump
  inflating: employeesdb/employees_partitioned.sql
  inflating: employeesdb/load_dept_manager.dump
  inflating: employeesdb/sql_test.sh
  inflating: employeesdb/load_departments.dump
  inflating: employeesdb/load_salaries1.dump
  inflating: employeesdb/show_elapsed.sql
  inflating: employeesdb/load_employees.dump
  inflating: employeesdb/sakila/README.md
  inflating: employeesdb/sakila/sakila-mv-data.sql
  inflating: employeesdb/sakila/sakila-mv-schema.sql
  inflating: employeesdb/images/employees.jpg
  inflating: employeesdb/images/employees.png
  inflating: employeesdb/images/employees.gif
theia@theiadocker-: /home/project$ cd employeesdb
theia@theiadocker-: /home/project/employeesdb$
```

4. In order to import the data, we'll need to load the data through MySQL. We can do that by navigating to the **Skills Network Toolbox**, selecting **Databases** and then selecting **MySQL**. Press **Start**. This will start a session of MySQL in SN Labs.




The **Inactive** label will change to **Starting**. This may take a few moments.

When it changes to **Active**, it means your session has started.



Take note of your password. You will need this to start MySQL.

You can copy your password by clicking the button next to it, as shown in the screenshot below:



5. With your password handy, we can now import the data. You can do this by entering the following into the Terminal:

```
1. mysql --host=127.0.0.1 --port=3306 --user=root --password -e < employees.sql
```

[Copy](#)

When prompted for your password, paste the password that you copied earlier into the Terminal and press **Enter**.

Please note, you won't be able to see your password when typing it in. Not to worry, this is expected!


6. Your data will now load. This may take a minute or so.

When you've finished loading the data, you'll see the following:



This means that your data has been imported.

7. To enter the MySQL command-line interface, return to your MySQL tab and select **MySQL CLI**.



8. Recall that the name of the database that we're using is **Employees**. To access it, we can use this command:

```
1. use employees
```

[Copy](#)

```
mysql> use employees
Reading table information for completion of table and column names
You can turn off this feature to get a quicker startup with -A

Database changed
```

9. Let's see which tables are available in this database:

```
1. show tables;
```

[Copy](#)


```
mysql> CREATE INDEX hire_date_index ON employees(hire_date);
Query OK, 0 rows affected (0.82 sec)
Records: 0 Duplicates: 0 Warnings: 0

mysql>
```

5. To check your index, you can use the `show index` command:

```
1. 1
2. SHOW INDEX FROM employees;
Copied
```

Now you can see that we have both the **emp_no** index and **hire_date** index.

```
mysql> SHOW INDEX FROM employees;
```

Table	Non_unique	Key_name	Seq_in_index	Column_name	Collation	Cardinality	Sub_part	Packed	Null	Index_type	Comment	Index_comment	Visible	Expression
employees	0	PRIMARY	1	emp_no	A	299423	NULL	NULL	NULL	BTREE			YES	NULL
employees	1	hire_date_index	1	hire_date	A	5324	NULL	NULL	NULL	BTREE			YES	NULL

2 rows in set (0.01 sec)

With the index added,

6. Once more, let's select all the employees who were hired on or after January 1, 2000.

```
1. 1
2. SELECT * FROM employees WHERE hire_date >= '2000-01-01';
Copied
```

```
mysql> SELECT * FROM employees WHERE hire_date >= '2000-01-01';
```

emp_no	birth_date	first_name	last_name	gender	hire_date
108201	1955-04-14	Mariangiola	Boreale	M	2000-01-01
60134	1964-04-21	Seshu	Rathonyi	F	2000-01-02
72329	1953-02-09	Randi	Luit	F	2000-01-02
424445	1953-04-27	Jeong	Boreale	M	2000-01-03
226533	1950-06-21	Xuejun	Benzmuller	F	2000-01-04
205040	1960-09-12	Ennio	Alblas	F	2000-01-06
227544	1954-11-17	Shahab	Demeyer	M	2000-01-08
422990	1953-04-09	Jaana	Verspoor	F	2000-01-11
47291	1960-09-09	Ulf	Flexer	M	2000-01-12
222965	1959-08-07	Volkmar	Perko	F	2000-01-13
499553	1954-05-06	Hideyuki	Delgrande	F	2000-01-22
428377	1957-05-09	Yucai	Gerlach	M	2000-01-23
463807	1964-06-12	Bikash	Covnot	M	2000-01-28

13 rows in set (0.00 sec)

```
mysql>
```

The difference is quite evident! Rather than taking about 0.17 seconds to execute the query, it takes 0.00 seconds—almost no time at all.

7. We can use the `EXPLAIN` statement to see how many rows were scanned:

```
1. 1
2. EXPLAIN SELECT * FROM employees WHERE hire_date >= '2000-01-01';
Copied
```

```
mysql> EXPLAIN SELECT * FROM employees WHERE hire_date >= '2000-01-01';
```

id	select_type	table	partitions	type	possible_keys	key	key_len	ref	rows	filtered	Extra
1	SIMPLE	employees	NULL	range	hire_date_index	hire_date_index	3	NULL	13	100.00	Using index condition

1 row in set, 1 warning (0.00 sec)

```
mysql>
```

Under **rows**, we can see that only the necessary 13 columns were scanned, leading to the improved performance.

Under **Extra**, you can also see that it has been explicitly stated that the index was used, that index being **hire_date_index** based on the **possible_keys** column.

Now, if you want to remove the index, enter the following into the Terminal:

```
1. 1
2. DROP INDEX hire_date_index ON employees;
Copied
```

This will remove the **hire_date_index** on the **employees** table. You can check with the `show index` command to confirm:

```
mysql> DROP INDEX hire_date_index ON employees;
Query OK, 0 rows affected (0.02 sec)
Records: 0 Duplicates: 0 Warnings: 0
```

```
mysql> SHOW INDEX FROM employees;
```

Table	Non_unique	Key_name	Seq_in_index	Column_name	Collation	Cardinality	Sub_part	Packed	Null	Index_type	Comment	Index_comment	Visible	Expression
employees	0	PRIMARY	1	emp_no	A	299423	NULL	NULL	NULL	BTREE			YES	NULL

1 row in set (0.00 sec)

Exercise 4: Use an UNION ALL Clause

Sometimes, you might want to run a query using the `OR` operator with `LIKE` statements. In this case, using a `UNION ALL` clause can improve the speed of your query, particularly if the columns on both sides of the `OR` operator are indexed.

1. To start, let's run this query:

```
1. 1
2. SELECT * FROM employees WHERE first_name LIKE 'C%' OR last_name LIKE 'C%';
Copied
```

emp_no	birth_date	first_name	last_name	gender	hire_date
499920	1953-07-18	Christ	Murtagh	M	1987-07-20
499933	1957-10-21	Chianti	Riesenhuber	F	1993-05-28
499936	1954-02-11	Chiranjit	Himler	M	1994-10-31
499947	1960-02-06	Conrado	Koyama	F	1989-02-19
499948	1953-05-24	Cordelia	Paludetto	M	1993-01-28
499956	1959-01-08	Zhonghua	Crooks	F	1994-10-12
499966	1955-12-04	Mihalis	Crabtree	F	1985-06-13
499975	1952-11-09	Masali	Chorvat	M	1992-01-23
499978	1960-03-29	Chiranjit	Kuzuoka	M	1990-05-24

28970 rows in set (0.20 sec)

This query searches for first names or last names that start with "C". It returned 28,970 rows, taking about 0.20 seconds.

2. Check using the `EXPLAIN` command to see how many rows are being scanned!

• Hint (Click Here)
• Solution (Click Here)

Once more, we can see that almost all the rows are being scanned, so let's add indexes to both the **first_name** and **last_name** columns.

3. Try adding an index to both the **first_name** and **last_name** columns.

• Hint (Click Here)
• Solution (Click Here)

4. Great! With your indexes now in place, we can re-run the query:

```
1. 1
2. SELECT * FROM employees WHERE first_name LIKE 'C%' OR last_name LIKE 'C%';
Copied
```

emp_no	birth_date	first_name	last_name	gender	hire_date
499881	1952-12-01	Christoph	Schneeberger	F	1987-10-29
499889	1956-01-29	Charlene	Hasham	F	1988-03-19
499908	1953-07-19	Toong	Coorg	F	1988-12-02
499916	1962-01-09	Florina	Cusworth	F	1997-05-18
499920	1953-07-18	Christ	Murtagh	F	1987-07-20
499933	1957-10-21	Chianti	Riesenhuber	F	1993-05-28
499936	1954-02-11	Chiranjit	Himler	M	1994-10-31
499947	1960-02-06	Conrado	Koyama	F	1989-02-19
499948	1953-05-24	Cordelia	Paludetto	M	1993-01-28
499956	1959-01-08	Zhonghua	Crooks	F	1994-10-12
499966	1955-12-04	Mihalis	Crabtree	F	1985-06-13
499975	1952-11-09	Masali	Chorvat	M	1992-01-23
499978	1960-03-29	Chiranjit	Kuzuoka	M	1990-05-24

28970 rows in set (0.16 sec)

Let's also see how many rows are being scanned:

```
1. 1
2. EXPLAIN SELECT * FROM employees WHERE first_name LIKE 'C%' OR last_name LIKE 'C%';
Copied
```

```
mysql> EXPLAIN SELECT * FROM employees WHERE first_name LIKE 'C%' OR last_name LIKE 'C%';
```

id	select_type	table	partitions	type	possible_keys	key	key_len	ref	rows	filtered	Extra
1	SIMPLE	employees	NULL	ALL	first_name_index,last_name_index	NULL	NULL	NULL	299423	20.99	Using where

1 row in set, 1 warning (0.00 sec)

With indexes, the query still scans all the rows.

5. Let's use the `UNION ALL` clause to improve the performance of this query.

We can do this with the following:

```
1. 1
1. SELECT * FROM employees WHERE first_name LIKE 'C%' UNION ALL SELECT * FROM employees WHERE last_name LIKE 'C%';
(Copied)
```

492481	1953-01-16	Chikara	Czap	M	1990-05-23
496850	1957-12-26	Cheong	Czap	F	1994-10-26

29730 rows in set (0.11 sec)

As we can see, this query only takes 0.11 seconds to execute, running faster than when we used the `IN` operator.

Using the `EXPLAIN` statement, we can see why that might be:

```
mysql> EXPLAIN SELECT * FROM employees WHERE first_name LIKE 'C%' UNION ALL SELECT * FROM employees WHERE last_name LIKE 'C%';
+----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| id | select_type | table | partitions | type | possible_keys | key | key_len | ref | rows | filtered | Extra |
+----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 1 | PRIMARY | employees | NULL | range | first_name_index | first_name_index | 58 | NULL | 20622 | 100.00 | Using index condition |
| 2 | UNION | employees | NULL | range | last_name_index | last_name_index | 66 | NULL | 34168 | 100.00 | Using index condition |
+----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
2 rows in set, 1 warning (0.00 sec)
```

As the `EXPLAIN` statement reveals, there were two select operations performed, with the total number of rows scanned sitting at 54,790. This is less than the original query that scanned the entire table and, as a result, the query performs faster.

Please note, if you choose to perform a leading wildcard search with an index, the entire table will still be scanned. You can see this yourself with the following query:

```
1. 1
1. SELECT * FROM employees WHERE first_name LIKE '%C';
(Copied)
```

With this query, we want to find all the employees whose first names end with "C".

When checking with the `EXPLAIN` and `SHOW INDEX` statements, we can see that although we have an index on `first_name`, the index is not used and results in a search of the entire table.

Under the `EXPLAIN` statement's **possible_keys** column, we can see that this index has not been used as the entry is `NULL`.

498090	1954-09-02	Marc	Fujisawa	F	1988-09-21
498599	1957-11-18	Marc	Awdeh	M	1986-07-25
499661	1963-06-30	Eric	Demeyer	M	1994-08-05

1180 rows in set (0.18 sec)

```
mysql> EXPLAIN SELECT * FROM employees WHERE first_name LIKE '%C';
+----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| id | select_type | table | partitions | type | possible_keys | key | key_len | ref | rows | filtered | Extra |
+----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 1 | SIMPLE | employees | NULL | ALL | NULL | NULL | NULL | NULL | 299423 | 11.11 | Using where |
+----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
1 row in set, 1 warning (0.00 sec)
```

```
mysql> SHOW INDEX from employees;
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| Table | Non_unique | Key_name | Seq_in_index | Column_name | Collation | Cardinality | Sub_part | Packed | Null | Index_type | Comment | Index_comment | Visible | Expression |
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| employees | 0 | PRIMARY | 1 | emp_no | A | 299423 | NULL | NULL | NULL | BTREE | | | YES | NULL |
| employees | 1 | first_name_index | 1 | first_name | A | 1251 | NULL | NULL | NULL | BTREE | | | YES | NULL |
| employees | 1 | last_name_index | 1 | last_name | A | 1585 | NULL | NULL | NULL | BTREE | | | YES | NULL |
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
3 rows in set (0.00 sec)
```

On the other hand, indexes do work with trailing wildcards, as seen with the following query that finds all employees whose first names begin with "C":

```
1. 1
1. SELECT * FROM employees WHERE first_name LIKE 'C%';
(Copied)
```

492080	1961-08-02	Cullen	Whittlesey	F	1997-01-12
495632	1958-05-16	Cullen	Pollock	M	1992-01-21

11294 rows in set (0.04 sec)

```
mysql> EXPLAIN SELECT * FROM employees WHERE first_name LIKE 'C%';
+----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| id | select_type | table | partitions | type | possible_keys | key | key_len | ref | rows | filtered | Extra |
+----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 1 | SIMPLE | employees | NULL | range | first_name_index | first_name_index | 58 | NULL | 20622 | 100.00 | Using index condition |
+----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
1 row in set, 1 warning (0.01 sec)
```

Under the `EXPLAIN` statement's **possible_keys** and **Extra** columns, we can see that the `first_name_index` is used. With only 20,622 rows scanned, the query performs better.

Exercise 5: Be SELECTive

In general, it's best practice to only select the columns that you need. For example, if you wanted to see the names and hire dates of the various employees, you could show that with the following query:

```
1. 1
1. SELECT * FROM employees;
(Copied)
```

499990	1956-09-05	Patricia	Breugel	M	1993-10-13
499999	1958-05-01	Sachin	Tsukuda	M	1997-11-30

300024 rows in set (0.26 sec)

```
mysql> EXPLAIN SELECT * FROM employees;
+----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| id | select_type | table | partitions | type | possible_keys | key | key_len | ref | rows | filtered | Extra |
+----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 1 | SIMPLE | employees | NULL | ALL | NULL | NULL | NULL | NULL | 299423 | 100.00 | NULL |
+----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
1 row in set, 1 warning (0.01 sec)
```

Notice how the query loads 300,024 rows in about 0.26 seconds. With the `EXPLAIN` statement, we can see that the entire table is being scanned, which makes sense because we are looking at all the entries.

If we, however, only wanted to see the names and hire dates, then we should select those columns:

```
1. 1
1. SELECT first_name, last_name, hire_date FROM employees;
(Copied)
```

Patricia	Breugel	1993-10-13
Sachin	Tsukuda	1997-11-30

300024 rows in set (0.17 sec)

```
mysql> EXPLAIN SELECT first_name, last_name, hire_date FROM employees;
+----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| id | select_type | table | partitions | type | possible_keys | key | key_len | ref | rows | filtered | Extra |
+----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 1 | SIMPLE | employees | NULL | ALL | NULL | NULL | NULL | NULL | 299423 | 100.00 | NULL |
+----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
1 row in set, 1 warning (0.00 sec)
```

As you can see, this query was executed a little faster despite scanning the entire table as well.

Give this a try!

Practice Exercise 1

Let's take a look at the `salaries` table. What if we wanted to see how much each employee earns?

When running the query, keep in mind how long it takes the query to run and how many rows are scanned each time.

1. First, let's select all the rows and columns from this table.

- Hint (Click Here)
- Solution (Click Here)

2. Now, let's see if there's a way to optimize this query. Since we only want to see how much each employee earns, then we can just select a few columns instead of all of them. Which ones would you select?

- Hint (Click Here)
- Solution (Click Here)

Practice Exercise 2

Let's take a look at the `department` table. What if we wanted to see the employee and their corresponding title?

Practice by selecting only the necessary columns and run the query!

- Hint (Click Here)
- Solution (Click Here)

Conclusion

Congratulations! Now, not only can you now identify common causes to slow queries, but you can resolve them by applying the knowledge that you have gained in this lab. Equipped with this problem-solving skill, you will be able to improve your queries performance, even in large databases.

Author(s)

Kathy An

Other Contributor(s)

Rav Abaja

Changelog

Date	Version	Changed by	Change Description
2021-10-05	1.0	Kathy An	Created initial version

© IBM Corporation 2021. All rights reserved.