

Estimated time needed: 45 minutes

In this lab, you'll learn how to monitor and optimize your database in PostgreSQL with both the command line interface (CLI) and database administration tool, pgAdmin.

Objectives

After completing this lab, you will be able to:

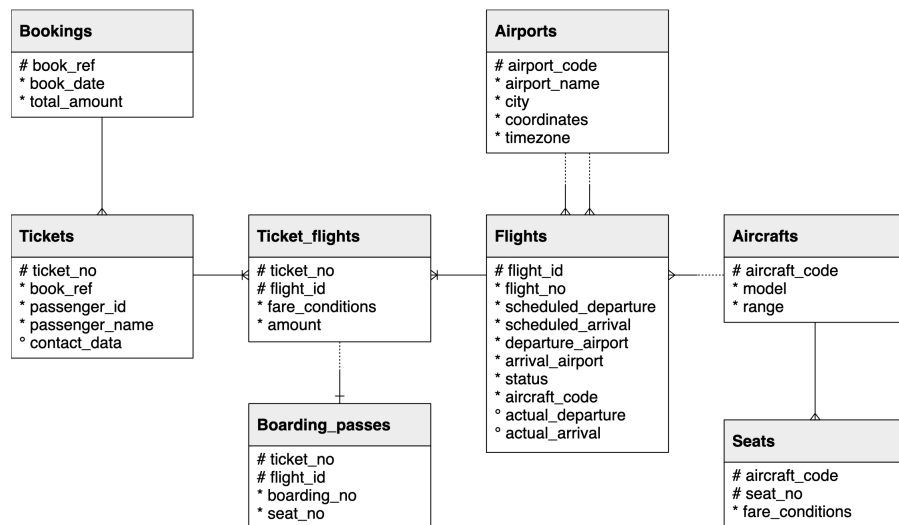
1. Monitor the performance of your database with the command line interface and pgAdmin.
2. Identify optimal data types for your database.
3. Optimize your database via the command line with best practices.

Software Used in this Lab

In this lab, you will be using PostgreSQL. It is a popular open source object relational database management system (RDBMS) capable of performing a wealth of database administration tasks, such as storing, manipulating, retrieving, and archiving data. To complete this lab, you will be accessing the PostgreSQL service through the IBM Skills Network (SN) Cloud IDE, which is a virtual development environment you will use throughout this course.

Database Used in this Lab

In this lab, you will use a database from <https://doctorespro.com/education/demosdb> distributed under the [PostgreSQL license](#). It stores a month of data about airline flights in Russia and is organized according to the following schema:

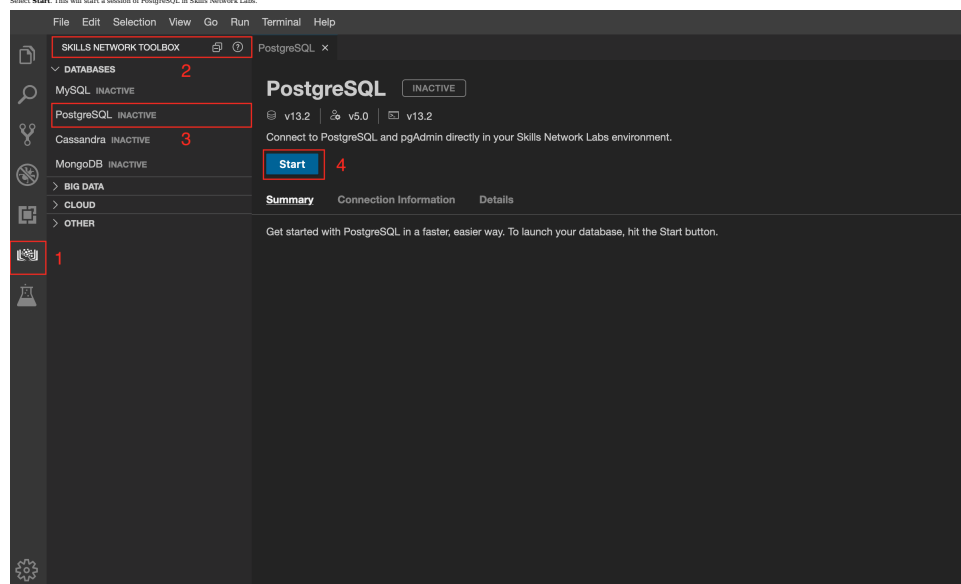


Exercise 1: Create Your Database

To get started with this lab, you'll launch PostgreSQL in Cloud IDE and create our database with the help of a SQL file.

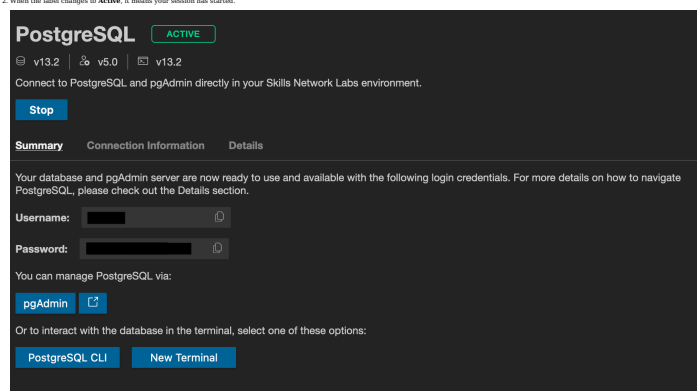
Task A: Start PostgreSQL in Cloud IDE

1. To start PostgreSQL, navigate to the **Skills Network Toolbox**, select **Databases**, and select **PostgreSQL**. Select **Start**. This will start a session of PostgreSQL in Skills Network Labs.



The **Inactive** label will change to **Starting**. This may take a minute or so.

2. When the label changes to **Active**, it means your session has started.



Task B: Create Your Database

1. Open a new terminal by selecting the **New Terminal** button near the bottom of the PostgreSQL tab

PostgreSQL

ACTIVE

v13.2

v5.0

v13.2

Connect to PostgreSQL and pgAdmin directly in your Skills Network Labs environment.

Stop

Summary

Connection Information

Details

Your database and pgAdmin server are now ready to use and available with the following login credentials. For more details on how to navigate PostgreSQL, please check out the Details section.

Username:

Password:

You can manage PostgreSQL via:

pgAdmin

PostgreSQL CLI

New Terminal

Or to interact with the database in the terminal, select one of these options:

2. With the terminal, you'll want to download the **demo** database that you're using in this lab. This database contains a smattering of data about flights in Russia.

To download it, you can use the following command:

```
1. 1
2. 2. wget https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/sample-guided-project/flights_RUSSIA_small.sql
3. 3. Copy
```

You should now see the SQL file in your file explorer in Cloud IDE.

File Edit Selection View Go Run Terminal Help

EXPLORER: PROJECT

postgres

flights_RUSSIA_small.sql

PostgreSQL

ACTIVE

v13.2

v5.0

v13.2

Connect to PostgreSQL and pgAdmin directly in your Skills Network Labs environment.

Stop

Summary

Connection Information

Details

Your database and pgAdmin server are now ready to use and available with the following login credentials. For more details on how to navigate PostgreSQL, please check out the Details section.

Username:

Password:

theia@theiadocker-: /home/project

theia@theiadocker-: /home/project\$ wget https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/example-guided-project/flights_RUSSIA_small.sql

--2021-10-13 16:04:19-- https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/example-guided-project/flights_RUSSIA_small.sql

Resolving cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud (cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud)... 198.23.119.245

Connecting to cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud (cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud)[198.23.119.245]:443... connected.

HTTP request sent, awaiting response... 200 OK

Length: 103865229 (99M) [application/x-sql]

Saving to: 'flights_RUSSIA_small.sql'

flights_RUSSIA_small.sql 100%[=====] 99.05M 30.8MB/s in 3.2s

--2021-10-13 16:04:22 (30.8 MB/s) - 'flights_RUSSIA_small.sql' saved [103865229/103865229]

theia@theiadocker-: /home/project\$

3. Let's return to the PostgreSQL tab and select the **PostgreSQL CLI** button near the bottom of the tab. This button will open a new PostgreSQL command line session.

PostgreSQL

ACTIVE

v13.2

v5.0

v13.2

Connect to PostgreSQL and pgAdmin directly in your Skills Network Labs environment.

Stop

Summary

Connection Information

Details

Your database and pgAdmin server are now ready to use and available with the following login credentials. For more details on how to navigate PostgreSQL, please check out the Details section.

Username:

Password:

You can manage PostgreSQL via:

pgAdmin

PostgreSQL CLI

New Terminal

Or to interact with the database in the terminal, select one of these options:

4. Now, you want to import the data from the file that you downloaded.

- Hint (Click Here)
- Solution (Click Here)

5. With our created database, let's see what tables you have. How many tables are there?

- Hint (Click Here)
- Solution (Click Here)

Great! With your environment and database set up, let's take a look at how you can monitor and optimize this database!

Exercise 2: Monitor Your Database

Database monitoring refers to reviewing the operational status of your database and maintaining its health and performance. With proper and proactive monitoring, databases will be able to maintain a consistent performance. Any problems that emerge, such as sudden outages, can be identified and resolved in a timely manner.

Tools such as pgAdmin, an open source graphical user interface (GUI) tool for PostgreSQL, come with several features that can help monitor your database. The main focus in this lab will be using the command line interface to monitor your database, but we'll also take a quick look at how the monitoring process can be replicated in pgAdmin.

Monitoring these statistics can be helpful in understanding your server and its databases, detecting any anomalies and problems that may arise.

Task A: Monitor Current Activity

To start, let's take a look at how you can monitor current server and database activity in PostgreSQL.

Server Activity

You can take a look at the server activity by running the following query:

```
1. 1
2. 2. SELECT pid, username, datname, state, state_change FROM pg_stat_activity;
```

Copy

This query will retrieve the following:

Column	Description
pid	Process ID
username	Name of user logged in
datname	Name of database
state	Current state, with two common values being: active (executing a query) and idle (waiting for new command)
state_change	Time when the state was last changed

This information comes from the **pg_stat_activity**, one of the built-in statistics provided by PostgreSQL.

1. Copy the query and paste it into the terminal.

You should see the following output:

```
demo=# SELECT pid, username, datname, state, state_change FROM pg_stat_activity;
 pid | username | datname | state | state_change
-----+-----+-----+-----+-----
 42  | postgres | postgres | idle  | 2021-10-13 22:11:20.330154+00
 51  | postgres | postgres | idle  | 2021-10-13 22:11:20.330154+00
1090 | postgres | demo    | active | 2021-10-13 22:11:20.725355+00
 40  |          |         |       |
 39  |          |         |       |
 41  |          |         |       |
(7 rows)
```

As you can see, there are currently 7 active connections to the server, with two of them being connected to databases that you're familiar with. After all, you started in the default **postgres** database, which is now idle, and now you're actively querying in the **demo** database.

2. To see what other columns are available for viewing, feel free to take a look at the [pg_stat_activity documentation](#). Let's say you wanted to see all the aforementioned columns, in addition to the actual text of the query that was last executed. Which columns should you add to review that?

- Hint (Click Here)
- Solution (Click Here)

Please note, if your table looks strange or squished, you can resize the terminal window by dragging it out.

If your result shows the text **(END)**, then type `q` to exit that view. Whenever you encounter this view, you can use `q` to return to your original view.

3. With queries, you can apply filtering. What if you only wanted to see the states that were **active**? How would you do that?

- Hint (Click Here)
- Solution (Click Here)

Database Activity

When looking at database activity, you can use the following query:

```
1. 1
1. SELECT datname, tup_inserted, tup_updated, tup_deleted FROM pg_stat_database;
```

`Copy`

This query will retrieve the following:

Column	Description
datname	Name of database
tup_inserted	Number of rows inserted by queries in this database
tup_updated	Number of rows updated by queries in this database
tup_deleted	Number of rows deleted by queries in this database

This information comes from the `pg_stat_database`, one of the statistics provided by PostgreSQL.

1. Copy the query and paste it into the terminal.

You should see the following output:

```
demo=# SELECT datname, tup_inserted, tup_updated, tup_deleted FROM pg_stat_database;
 datname | tup_inserted | tup_updated | tup_deleted 
-----+-----+-----+-----
 postgres |          2  |          1 |          0 
 demo    | 2290162    |         22 |          0 
 template1 |          0  |          0 |          0 
 template0 |          0  |          0 |          0 
(5 rows)
```

As you can see, the two databases that are returned are the `postgres` and `demo`. These are databases that you are familiar with. The other two, `template1` and `template0` are default templates for databases, and can be overlooked in this analysis. Based on this output, you now know that `demo` had about 2,290,162 rows inserted and 22 rows updated.

2. To see what other columns are available for viewing, you can read through the [pg_stat_database documentation](#).

Let's say you wanted to see the number of rows fetched and returned by this database.

Note: The number of rows fetched is the number of rows that were returned. The number of rows returned is the number of rows that were read and scanned by the query.

What query should you use to do that?

- Hint (Click Here)
- Solution (Click Here)

3. With queries, you can apply filtering. What if you only wanted to see the database details (rows inserted, updated, deleted, returned and fetched) for `demo`?

- Hint (Click Here)
- Solution (Click Here)

Later, we'll take a look at how you can monitor these activities in pgAdmin.

Task B: Monitor Performance Over Time

Extensions, which can enhance your PostgreSQL experience, can be helpful in monitoring your database. One such extension is `pg_stat_statements`, which gives you an aggregated view of query statistics.

1. To enable the extension, enter the following command:

```
1. 1
1. CREATE EXTENSION pg_stat_statements;
```

`Copy`

This will enable the `pg_stat_statements` extension, which will start to track the statistics for your database.

2. Now, let's edit the PostgreSQL configuration file to include the extension you just added:

```
1. 1
1. ALTER SYSTEM SET shared_preload_libraries = 'pg_stat_statements';
```

`Copy`

For the changes to take effect, you will have to restart your database. You can do that by typing `exit` in the terminal to stop your current session. Close the terminal and return to the PostgreSQL tab. Select **Stop**.

PostgreSQL

ACTIVE

v13.2 | v5.0 | v13.2

Connect to PostgreSQL and pgAdmin directly in your Skills Network Labs environment.

Stop

Summary | Connection Information | Details

Your database and pgAdmin server are now ready to use and available with the following login credentials. For more details on how to navigate PostgreSQL, please check out the Details section.

Username:

Password:

You can manage PostgreSQL via:

pgAdmin

Or to interact with the database in the terminal, select one of these options:

PostgreSQL CLI | New Terminal

When the session has become **Inactive** once more, select **Start** to restart your session.

3. Once your session has started, open the **PostgreSQL CLI**.

You'll need to reconnect to the `demo` database, which you can do by using the following command:

```
1. 1
1. \connect demo;
```

`Copy`

4. You can see if this extension has been loaded by checking both the installed extensions and the `shared_preload_libraries`.

First let's check the installed extensions:

```
1. 1
1. \l
```

`Copy`

Name	Version	Schema	List of installed extensions	Description
pg_stat_statements	1.5.0	postgres		Stats planning and execution statistics of all SQL statements executed by users
pg_catalog	1.5.0	pg_catalog	Plpgsql, procedural, language	

Notice how `pg_stat_statements` has been installed.

You can also check the `shared_preload_libraries` with:

```
1. 1
1. \show shared_preload_libraries;
```

`Copy`

```
demo=# \show shared_preload_libraries;
shared_preload_libraries
-----
pg_stat_statements,
libpq
```

`pg_stat_statements` is also shown under `shared_preload_libraries`.

5. Since the results returned by `pg_stat_statements` can be quite long, let's turn on expanded table formatting with the following command:

```
1. 1
1. \x
```

`Copy`

This will display the output tables in an expanded table format.

```
demo=# \x
Expanded display is on.
demo=#
```

You can turn it off by repeating the `\x` command.

6. From the [pg_stat_statements documentation](#), you'll see the various columns available to be retrieved.

Let's say you wanted to retrieve the database ID, the query, and total time that it took to execute the statement (in milliseconds).

- Hint (Click Here)
- Solution (Click Here)

7. What if you wanted to check which database name matches the database ID?

- Hint (Click Here)
- Solution (Click Here)

It's important to note that adding these extensions can increase your server load, which may affect performance. If you need to drop the extension, you can achieve that with the following command:

```
1. 1
1. DROP EXTENSION pg_stat_statements;
```

`Copy`

If you check the current extensions with `\x`, you'll also see that `pg_stat_statements` no longer appears.

You should also reset the `shared_preload_libraries` in the configuration file:

```
1. 1
1. ALTER SYSTEM RESET shared_preload_libraries;
```

`Copy`

After this, you'll need to exit the terminal and restart the PostgreSQL CLI to see the changes reflected in `\show shared_preload_libraries`.

Task C: Monitor with pgAdmin

Another method of monitoring your database comes in the form of pgAdmin. In order to use this tool, you'll need to first launch it.

1. Open the PostgreSQL tab from the Skills Network Tooltop and select the pop-out button next to the **pgAdmin** button. This will open pgAdmin in a new tab.

PostgreSQL

ACTIVE

v13.2 | v5.0 | v13.2

Connect to PostgreSQL and pgAdmin directly in your Skills Network Labs environment.

Stop

Summary | Connection Information | Details

Your database and pgAdmin server are now ready to use and available with the following login credentials. For more details on how to navigate PostgreSQL, please check out the Details section.

Username:

Password:

You can manage PostgreSQL via:

pgAdmin

Or to interact with the database in the terminal, select one of these options:

PostgreSQL CLI

New Terminal

2. In the left panel, select the dropdown next to **Servers**. You'll be prompted to enter a password.

pgAdmin

Management

Feature rich | Maximise

pgAdmin is an Open Source admin

is designed to answer the needs of

Connect to Server

Please enter the password for the user 'postgres' to connect the server - 'postgres'

Password

☐ Save Password

Cancel OK

Welcome

Dashboard | Properties | SQL | Statistics | Dependencies | Dependents

Browser

Servers (1)

1 > postgres

Quick Links

Add New Server

Configure pgAdmin

Getting Started

PostgreSQL Documentation

pgAdmin Website

Planet PostgreSQL

Community Support

3. Return to your Cloud IDE session. In the PostgreSQL tab, select the copy button next to the **Password** field. This is the password that you can enter into pgAdmin.

PostgreSQL

ACTIVE

v13.2 | v5.0 | v13.2

Connect to PostgreSQL and pgAdmin directly in your Skills Network Labs environment.

Stop

Summary | Connection Information | Details

Your database and pgAdmin server are now ready to use and available with the following login credentials. For more details on how to navigate PostgreSQL, please check out the Details section.

Username:

Password:

You can manage PostgreSQL via:

pgAdmin

Or to interact with the database in the terminal, select one of these options:

PostgreSQL CLI

New Terminal

Paste that password into pgAdmin. Then, click **OK**.

Your server will now load.

4. On the home page, under **Dashboard**, you will have access to server or database statistics, depending on which you are looking at.

Dashboard | Properties | SQL | Statistics | Dependencies | Dependents

Server sessions

Transactions per second

Tuples in

Tuples out

Block I/O

Server activity

Sessions | Locks | Prepared Transactions | Configuration

Search

	PID	Database	User	Application	Client	Backend start	State	Wait event	Blocking PIDs
✖	40					2021-10-14 02:05:01 UTC		Activity: CheckpointerMain	
✖	41					2021-10-14 02:05:01 UTC		Activity: BgWriterMain	
✖	42					2021-10-14 02:05:01 UTC		Activity: WalWriterMain	
✖	43					2021-10-14 02:05:01 UTC		Activity: AutoVacuumMain	
✖	45		postgres			2021-10-14 02:05:01 UTC		Activity: LogicalLauncherMain	
✖	52	postgres	postgres	sysdig-agent		2021-10-14 02:05:09 UTC	idle	Client: ClientRead	
✖	271	postgres	postgres	pgAdmin 4 - DB:postgres	172.18.0.3	2021-10-14 02:08:07 UTC	active		

The table below lists the displayed statistics on the Dashboard that correspond with the statistics that you accessed with the CLI.

Chart Description
Server/Database sessions Displays the total sessions that are running. For servers, this is similar to the `pg_stat_activity`, and for databases, this is similar to the `pg_stat_database`.
Transactions per second Displays the commits, rollbacks, and transactions taking place.
Tuples in Displays the number of tuples (rows) that have been inserted, updated, and deleted, similar to the `tup_inserted`, `tup_updated`, and `tup_deleted` columns from `pg_stat_database`.
Tuples out Displays the number of tuples (rows) that have been fetched (returned as output) or returned (read or scanned). This is similar to `tup_fetched` and `tup_returned` from `pg_stat_database`.

- Chart** Displays the sessions, locks, prepared transactions, and configuration for the server. In the Sessions tab, it offers a look at the breakdown of the sessions that are currently active on the server, similar to the view provided by `pg_stat_activity`. To check for any new processes, you can select the refresh button at the top-right corner.
- Description**
5. You can test these charts out by starting another session.
- Return to the tab with the Cloud IDE environment. On the PostgreSQL tab, select **PostgreSQL CLI**. This will start a new session of PostgreSQL with the CLI.

PostgreSQL

ACTIVE

v13.2 | v5.0 | v13.2

Connect to PostgreSQL and pgAdmin directly in your Skills Network Labs environment.

Stop

Summary Connection Information Details

Your database and pgAdmin server are now ready to use and available with the following login credentials. For more details on how to navigate PostgreSQL, please check out the Details section.

Username:

Password:

You can manage PostgreSQL via:

pgAdmin

Or to interact with the database in the terminal, select one of these options:

PostgreSQL CLI **New Terminal**

6. Once you have started that instance, switch back to the tab with pgAdmin.

- What do you notice?
- Hint (Click Here)
 - Solution (Click Here)

7. To see the dashboard for the **demo** database, navigate to the left panel and select the **Databases** dropdown and then select the **demo** database to connect to it. As you can see, similar statistics are displayed for the database.

pgAdmin

File Object Tools Help

Browser Servers (1) postgres Databases (2) demo Casts Catalogs Event Triggers Extensions Foreign Data Wrappers Languages Publications Schemas Subscriptions postgres Login/Group Roles Tablespaces

Dashboard Properties SQL Statistics Dependencies Dependents

Database sessions Total Active Idle

Transactions per second Transactions Commits Rollbacks

Tuples in Inserts Updates Delete

Tuples out Fetched Returned

Block I/O Reads Hits

Server activity Sessions Locks Prepared Transactions

		PID	User	Application	Client	Backend start	State	Wait event	Blocking PIDs
+	■	951	postgres	pgAdmin 4 - DB:demo	172.18.0.3	2021-10-14 02:17:21 UTC	active		

8. Let's run a query on the database! To do that, navigate to the menu bar and select **Tools** > **Query Tool**.

You can run any query. To keep things simple, let's run the following to select all the data from the **bookings** table:

```
1. SELECT * FROM bookings;
```

Select the run button. You will see that the query has successfully loaded.

pgAdmin

demo/postgres@postgres *

Query Editor Query History

```
1 SELECT * FROM bookings;
```

Data Output Explain Messages Notifications

	book_ref [PK] character (5)	book_date timestamp with time zone	total_amount numeric (10,2)
1	00000F	2017-07-05 00:12:00+00	265700.00
2	000012	2017-07-14 06:02:00+00	37900.00
3	000068	2017-08-15 11:27:00+00	18100.00
4	000181	2017-08-10 10:28:00+00	131800.00
5	0002D8	2017-08-07 18:40:00+00	23600.00
6	0002DB	2017-07-29 03:30:00+00	101500.00
7	0002E0	2017-07-11 13:09:00+00	89600.00
8	0002F3	2017-07-10 02:31:00+00	69600.00
9	00034E	2017-08-04 13:52:00+00	73300.00
10	000352	2017-07-05 23:02:00+00	109500.00
11	000374	2017-08-12 07:13:00+00	136200.00

Successfully run. Total query runtime: 604 msec. 262788 rows affected.

9. In pgAdmin, switch back to the database's **Dashboard** tab. You can refresh the **Server activity** and check to see if any of the charts have shown a spike since the data was retrieved.

- What do you notice?
- Hint (Click Here)
 - Solution (Click Here)

While you can monitor your database through the command line alone, tools like pgAdmin can be helpful in providing a visual representation of how your server and its databases are performing. PostgreSQL also offers logging capabilities to monitor and troubleshoot your lab, which will be further discussed in the Troubleshooting lab.

Exercise 3: Optimize Your Database

Data optimization is the maximization of the speed and efficiency of retrieving data from your database. Optimizing your database will improve its performance, whether that's inserting or retrieving data from your database. Doing this will improve the experience of anyone interacting with the database. Similar to MySQL, there are optimal data types and maintenance (otherwise known as "vacuuming") that can be applied to optimize databases.

Task A: Optimize Data Types

When it comes to optimizing data types, understanding the data values will help in selecting the proper data type for the column.

Let's take a look at an example in the **demo** database.

- Return to the CLI session that you opened previously (or open a new session if it has been closed).

- [Hint \(Click Here\)](#)
- [Solution \(Click Here\)](#)

2. 2
2. 2
Copied!

- [Hint \(Click Here\)](#)
- [Solution \(Click Here\)](#)

- [Hint \(Click Here\)](#)
- [Solution \(Click Here\)](#)

```
1. 1
1. \d aircraft_data
```

2. 2
2. DROP VOW aircraft;

1. 1
1. ALTER TABLE aircraft

5. Now, let's check the table's columns and data types again:

• [Hint \(Click Here\)](#)

in your day-to-day life, you can vacuum our rooms to keep them neat and tidy. You can do the same with databases by maintaining and optimizing them with some vacuuming

¹ In PostgreSQL, vacuuming means to clean out your databases by reclaiming any storage from "dead tuples", otherwise known as rows that have been deleted but have not been cleaned out

Generally, the **autovacuum** feature is automatically enabled, meaning that PostgreSQL will automate the vacuum maintenance process for you.

1. 1
1. show answers

As you can see, **autovacuum** is enabled.

Since `autovacuum` is enabled, let's check to see when your database was last vacuumed.

to do that, you can use the `pg_stat_user_tables`, which displays statistics about each table that is a user table (instead of a system table) in the database. The columns that are returned from `pg_stat_user_tables` are:

[Hint \(Click Here\)](#)

congratulations! Now, not only do you know how to monitor and optimize your database with the CLI, but you can also do so with `netAdmin`. You will now be able to apply this knowledge to any PostgreSQL databases you create and modify in the future.

 $\text{Leptothorax}(n)$

Author(s)

Lathy An

Changelog

Date	Version	Changed by	Change Description
2021-10-14	1.0	Kathy An	Created initial version

© IBM Corporation 2021. All rights reserved.